

Predicting Marketing Campaign Response Using Logistic Regression

by Rabindranatah Duran Pons, Yasaman Eftekharypour, Valeria Siciliano, Rocco Lee

Summary

In this project, we developed a machine-learning pipeline using logistic regression to predict whether a customer will subscribe to a marketing campaign. The workflow combined a preprocessing stage (StandardScaler and OneHotEncoder) with a logistic regression classifier, followed by training and evaluation using a train/test split.

After applying class-weighting to address the dataset's imbalance, the model achieved an accuracy of approximately 85% and a ROC-AUC of ~0.91, indicating strong overall discrimination between the subscribed ("yes") and not-subscribed ("no") classes. Importantly, class-weighting significantly improved the model's ability to detect positive cases, giving the "yes" class a recall of 0.81. This shows that the weighted logistic regression approach is better suited for imbalanced marketing data, where correctly identifying potential subscribers is more valuable than simply maximizing accuracy.

Introduction

Background:

Financial institutions frequently rely on direct marketing campaigns to promote new products and services to potential clients. One common approach is telephone-based marketing, where bank representatives call customers to inform them about financial offerings, such as long-term deposits and attempt to persuade them to subscribe. Although this method can be effective, it is also resource-intensive calling uninterested or unlikely customers wastes time, labor, and operational costs. As a result, banks increasingly turn to data-driven decision-making to identify which individuals are most likely to respond positively to a campaign.

The Bank Marketing Dataset was created within this context. Compiled by researchers at the University of Minho in Portugal, it captures detailed information from a series of telephone marketing campaigns run by a Portuguese banking institution. The dataset includes demographic attributes (such as age, employment type, and marital status), financial indicators (such as credit defaults and loan status), and campaign-specific variables (including previous contact outcomes and call duration). The target variable

indicates whether a customer ultimately subscribed to a term deposit, making the dataset a classic example of a binary classification problem.

The main question explored in this project is:

"Can we predict whether a bank customer will subscribe to a term deposit based on their demographic characteristics, financial information, and interactions with previous marketing campaigns?"

The dataset includes three main categories of features:

1. Client demographics and personal information

- age
- job
- marital
- education
- default (has credit in default)
- housing (has housing loan)
- loan (has personal loan)

2. Current campaign interaction

- contact — type of communication (cellular/telephone)
- day_of_week — day of contact
- month — month of campaign
- duration — call duration in seconds
- campaign — number of contacts during this campaign

3. Past campaign and historical interaction

- pdays — number of days since last contact
- previous — number of previous contacts
- poutcome — outcome of previous campaign

Methods

Data

The dataset used in this project is the Bank Marketing Dataset, created by Moro, Cortez, and Rita (2014) at the University of Minho in Portugal. The data was sourced from the UCI Machine Learning Repository, and can be accessed online at <https://archive.ics.uci.edu/dataset/222/bank+marketing>. Each row in the dataset represents a single customer contacted during a direct marketing phone campaign, and includes information such as demographic attributes, financial status, call details, previous campaign interactions, and the final outcome indicating whether the customer subscribed to a term deposit.

Analysis

A logistic regression model was used to predict whether a marketing campaign will be successful or not. All original variables from the dataset were included in the analysis. Before fitting, numerical features were standardized with a StandardScaler, and categorical variables were converted to binary indicators via OneHotEncoder. The dataset was split into 80% training and 20% testing, and class imbalance was addressed by balancing class weights during model training. The model's performance was evaluated using accuracy and ROC-AUC scores.

The code used to perform this analysis and generate the accompanying report can be found here:

https://github.com/Roccolee18/bank_marketing_group_24/blob/main/marketing_campaign_pipeline.ipynb

Results & Discussion

The logistic regression model developed for this analysis provides meaningful insight into the factors associated with customer subscription, but it also highlights the intrinsic challenges of modeling imbalanced marketing data. Our pipeline combined appropriate preprocessing steps—StandardScaler for numerical features and One-Hot Encoding for categorical variables—with a LogisticRegression classifier to ensure proper handling of the heterogeneous dataset while respecting the Golden Rule and avoiding data leakage.

The performance metrics indicate that the model performs reasonably well overall. The ROC-AUC score of approximately 0.91 suggests strong ability to distinguish between subscribers ("yes") and non-subscribers ("no"). Although overall accuracy is around 0.85, accuracy alone is not an appropriate metric for this imbalanced context, because the majority class dominates the dataset.

More importantly, the class-weighted logistic regression successfully shifts the model's focus toward the minority class. The recall for the "yes" class reaches 0.81, a substantial

improvement compared to what a non-weighted model would typically achieve on an imbalanced dataset. This indicates that the model is able to identify most customers who eventually subscribe—an outcome that aligns with the core business objective, where failing to detect potential subscribers is far more costly than incorrectly flagging non-subscribers. The precision for the “yes” class is lower (0.42), which is an expected trade-off: by increasing recall and giving more weight to positive cases, the classifier becomes more permissive and produces more false positives. However, in a marketing context—where the cost of contacting an uninterested customer is low compared to the value of identifying a true potential subscriber—this trade-off is acceptable and strategically desirable.

The confusion matrix supports this interpretation. Out of 1,058 actual subscribers, the model correctly identifies 862 true positives while misclassifying 196 as non-subscribers. On the other hand, among the majority class, 6,786 non-subscribers are correctly classified, with 1,199 false positives. These numbers reflect a deliberate shift in the decision boundary due to class balancing: the model becomes more sensitive to the minority class at the expense of increasing false positives.

Overall, the balanced logistic regression model is appropriate for this business problem. Its ability to capture a large portion of true subscribers, even with lower precision, aligns with the strategic goal of maximizing successful marketing outreach. By prioritizing recall in the positive class, the model supports proactive customer engagement and provides a meaningful foundation for future marketing campaigns.

These results show that using a class-weighted logistic regression helps the model catch many more people who are likely to subscribe. This can be useful for marketing teams because it means they can focus their efforts on customers who are more likely to say “yes.” It also shows which factors—like the success of previous campaigns, the month of contact, or call duration—matter most, which can help improve how future campaigns are planned.

These results also bring up a number of future questions. For example, it is unclear whether another type of model, such as a tree-based method, could perform even better than logistic regression on this imbalanced data. Another question is whether the same patterns would appear if we ran this analysis on a different marketing campaign or a different time period. Finally, it would be useful to understand which types of customers the model tends to misclassify most often, and whether adding more customer information could help the model make more reliable predictions.

The following code reads the data programatically and saves it to the data folder:

```
In [1]: from ucimlrepo import fetch_ucirepo  
import pandas as pd  
import altair as alt  
import os
```

```

import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.compose import ColumnTransformer
from sklearn.pipeline import Pipeline
from sklearn.metrics import (
    accuracy_score, classification_report,
    roc_auc_score, confusion_matrix
)
from sklearn.linear_model import LogisticRegression
import matplotlib.pyplot as plt
import seaborn as sns

# Create the data folder if it doesn't exist
os.makedirs("data", exist_ok=True)

# fetch dataset
bank_marketing = fetch_ucirepo(id=222)

# Convert the data into a pd dataframe
X = bank_marketing.data.features
y = bank_marketing.data.targets

df = pd.concat([X, y], axis=1)

# Save combined dataset to data folder
df.to_csv("data/bank_marketing.csv", index=False)

print(df.shape)
print(df.head())

```

(45211, 17)

	age	job	marital	education	default	balance	housing	loan	\
0	58	management	married	tertiary	no	2143	yes	no	
1	44	technician	single	secondary	no	29	yes	no	
2	33	entrepreneur	married	secondary	no	2	yes	yes	
3	47	blue-collar	married	NaN	no	1506	yes	no	
4	33	NaN	single	NaN	no	1	no	no	

	contact	day_of_week	month	duration	campaign	pdays	previous	poutcome	y
0	NaN	5	may	261	1	-1	0	NaN	no
1	NaN	5	may	151	1	-1	0	NaN	no
2	NaN	5	may	76	1	-1	0	NaN	no
3	NaN	5	may	92	1	-1	0	NaN	no
4	NaN	5	may	198	1	-1	0	NaN	no

DATA VALIDATION

Before working on the dataset, we need to perform some validation checks to make sure that raw data meet the required format. The validation will be performed on the following points:

- Correct data file format
- Correct column names
- No empty observations
- Missingness not beyond expected threshold
- Correct data types in each column
- No duplicate observations
- No outlier or anomalous values
- Correct category levels (i.e., no string mismatches or single values)
- Target/response variable follows expected distribution
- No anomalous correlations between target/response variable and features/explanatory variables
- No anomalous correlations between features/explanatory variables

Correct data file format

```
In [2]: from pathlib import Path
import pandera as pa
from pandera import Column, DataFrameSchema

def validate_correct_data_format(df, file_path):
    """
    Validate that the data format matches .csv.

    Parameters
    -----
    df : pandas.DataFrame
        The used dataset.
    file_path : str or Path
        Path to the raw data file.

    """
    file_path = Path(file_path)

    if file_path.suffix.lower() != ".csv":
        raise ValueError(
            f"File format is not correct: expected '.csv', got '{file_path.suffix}'"
        )
```

```
)  
  
print("The file format validation passed: .csv file detected")  
  
validate_correct_data_format(df, "data/bank_marketing.csv")
```

The file format validation passed: .csv file detected

Correct column names

```
In [3]: from pandera import Column, DataFrameSchema  
  
columns = [  
    "age", "job", "marital", "education", "default",  
    "balance", "housing", "loan", "contact",  
    "day_of_week", "month", "duration", "campaign",  
    "pdays", "previous", "poutcome", "y"  
]  
column_schema=DataFrameSchema(  
    columns={c: Column(nullable=True) for c in columns},  
    strict=True,  
    ordered=False  
)  
def column_names(df):  
    """  
    Validate of the correct column names.  
    """  
  
    column_schema.validate(df)  
    print("Correct column names were passed")  
  
column_names(df)
```

Correct column names were passed

No empty observations

```
In [4]: def no_empty_observations(df):  
    """  
    Validation to make sure there are no empty observations (row where all v  
    """  
    n_obs=df.isna().all(axis=1).sum()  
  
    empty_obs=pa.DataFrameSchema(  
        columns=None,  
        checks=[  
            pa.Check(  
                lambda df: ~(df.isna().all(axis=1)).any(),  
                element_wise=False,  
                error="Found one or more completely empty rows in the dataset."  
        ]  
    )  
    empty_obs.validate(df)  
    print(f"{n_obs} empty observations found in dataset.")
```

```
no_empty_observations(df)
```

```
0 empty observations found in dataset.
```

No anomalous correlations between target/response variable and features/explanatory variables

```
In [5]: from deepchecks.tabular import Dataset
from deepchecks.tabular.checks import FeatureLabelCorrelation

bank_train_ds = Dataset(
    df, label="y",
    cat_features=['job',
                  'marital',
                  'education',
                  'default',
                  'housing',
                  'loan',
                  'contact',
                  'month',
                  'poutcome']
)

# Trigger a warning if any feature has a predictive power score (correlation
check_feat_lab_corr = FeatureLabelCorrelation().add_condition_feature_pps_le
check_feat_lab_corr_result = check_feat_lab_corr.run(dataset=bank_train_ds)

if not check_feat_lab_corr_result.passed_conditions():
    raise ValueError("Feature-Label correlation exceeds the maximum acceptab
```

```
/Users/rabin/miniforge3/envs/522/lib/python3.12/site-packages/deepchecks/cor
e/serialization/dataframe/html.py:16: UserWarning:
```

```
pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg\_resources.html. The pkg_resources package is slated for removal as early as 2025-11-30. Refrain from using this package or pin to Setuptools<81.
```

No anomalous correlations between features/explanatory variables

```
In [6]: from deepchecks.tabular.checks import FeatureFeatureCorrelation

check_feat_feat_corr = FeatureFeatureCorrelation().add_condition_max_number_
    threshold=0.99,
)

# Run the check
check_feat_feat_corr_result = check_feat_feat_corr.run(dataset=bank_train_ds

# Uncomment to print the actual values
#print(check_feat_feat_corr_result.value)
```

```
# Trigger a warning if features have a predictive power (correlation) greater
if not check_feat_feat_corr_result.passed_conditions():
    raise ValueError("Feature-feature correlation exceeds the maximum accept
```

EDA

The following code performs some preliminary EDA:

In [7]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 45211 entries, 0 to 45210
Data columns (total 17 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   age          45211 non-null   int64  
 1   job           44923 non-null   object  
 2   marital       45211 non-null   object  
 3   education     43354 non-null   object  
 4   default       45211 non-null   object  
 5   balance        45211 non-null   int64  
 6   housing        45211 non-null   object  
 7   loan           45211 non-null   object  
 8   contact        32191 non-null   object  
 9   day_of_week    45211 non-null   int64  
 10  month          45211 non-null   object  
 11  duration       45211 non-null   int64  
 12  campaign       45211 non-null   int64  
 13  pdays          45211 non-null   int64  
 14  previous       45211 non-null   int64  
 15  poutcome       8252 non-null   object  
 16  y               45211 non-null   object  
dtypes: int64(7), object(10)
memory usage: 5.9+ MB
```

In [8]: `df.describe()`

	age	balance	day_of_week	duration	campaign	previous
count	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000	45211.000000
mean	40.936210	1362.272058	15.806419	258.163080	2.763841	4.000000
std	10.618762	3044.765829	8.322476	257.527812	3.098021	10.000000
min	18.000000	-8019.000000	1.000000	0.000000	1.000000	-1.000000
25%	33.000000	72.000000	8.000000	103.000000	1.000000	-1.000000
50%	39.000000	448.000000	16.000000	180.000000	2.000000	-1.000000
75%	48.000000	1428.000000	21.000000	319.000000	3.000000	-1.000000
max	95.000000	102127.000000	31.000000	4918.000000	63.000000	87.000000


```
In [9]: df.value_counts()
```

```
Out[9]: age      job      marital      education      default      balance      housing      loan      con
tact      day_of_week      month      duration      campaign      pdays      previous      poutcome
y
18      student      single      primary      no      608      no      no      cel
lular      13      nov      210      1      93      1      success
yes      1
44      blue-collar      married      secondary      no      6491      yes      no      cel
lular      19      nov      126      2      174      1      failure
no      1
lular      20      apr      87      2      3060      yes      no      cel
no      1      failure
lular      18      may      18      5      2979      yes      no      cel
no      1      other
lular      1      oct      82      1      1495      yes      no      cel
no      1      failure
..
34      management      single      tertiary      no      -444      yes      yes      cel
lular      17      apr      129      2      148      4      other
no      1
lular      12      married      tertiary      no      8000      no      no      cel
yes      1      feb      291      1      260      2      failure
lular      29      jan      159      7      5878      no      no      cel
no      1      success
lular      30      apr      113      1      4859      no      no      cel
no      1
89      retired      divorced      primary      no      1323      no      no      tel
ephone      29      dec      207      4      189      1      other
no      1
Name: count, Length: 7842, dtype: int64
```

The distribution of people who subscribed and didn't subscribe can be found below:


```
In [10]: # Simplify working with large datasets in Altair
alt.data_transformers.enable('vegafusion')

alt.Chart(df).mark_bar().encode(
    y=alt.Y("y:N", title="Subscribed"),
    x=alt.X("count()", title="Count"),
    color="y:N"
)
```



A comparison of the distribution of subscribed people among age can be found below:

```
In [11]: (
    alt.Chart(df)
        .transform_density(
            density="age",
            groupby=["y"],
            as_=["age", "density"]
        )
        .mark_area(opacity=0.4)
        .encode(
            x=alt.X("age:Q", title="Age"),
            y=alt.Y("density:Q", title="Density"),
            color=alt.Color("y:N", title="Subscribed")
        )
)
```



This shows that age can be a very good predictor for subscription when used on its own.


```
In [12]: (
    alt.Chart(df)
        .transform_density(
            density="balance",
            groupby=["y"],
            as_=["balance", "density"]
        )
        .mark_area(opacity=0.4)
```

```

.encode(
    x=alt.X("balance:Q", title="Balance"),
    y=alt.Y("density:Q", title="Density"),
    color=alt.Color("y:N", title="Subscribed")
)
.properties(title="Balance Density by Subscription Outcome")
)

```

Out[12]:



In the plot above we can see that the distribution of balance is right skewed and it might not be a good predictor on its own for subscription. However, when considering other variables, it might be useful.

The following shows a heatmap of the correlation among all the variables:


```

In [13]: numeric_cols = df.select_dtypes(include='number').columns
corr_df = df[numeric_cols].corr().stack().reset_index()
corr_df.columns = ['var1', 'var2', 'correlation']

alt.Chart(corr_df).mark_rect().encode(
    x=alt.X('var1:N', title=""),
    y=alt.Y('var2:N', title=""),
    color=alt.Color('correlation:Q', scale=alt.Scale(scheme='redblue')),
    tooltip=['var1', 'var2', 'correlation']
).properties(
    title="Correlation Matrix"
)

```

Out[13]:



We can see that there is no multicollinearity among different variables.

```
In [14]: print(df['education'].unique())
print(df['marital'].unique())
```

```
['tertiary' 'secondary' nan 'primary']
['married' 'single' 'divorced']
```

Preprocessing

```
In [15]: df = df.dropna()
df = df[df['education'] != 'unknown']
df = df[df['job'] != 'unknown']
df = df[df['marital'] != 'unknown']

print(df.head())
print(df.info())
```

```

      age      job marital education default balance housing loan \
24060  33    admin. married   tertiary    no     882     no  no
24062  42    admin. single secondary   no    -247    yes yes
24064  33    services married secondary   no    3444    yes  no
24072  36 management married tertiary   no    2415    yes  no
24077  36 management married tertiary   no       0    yes  no

      contact day_of_week month duration campaign pdays previous \
24060 telephone        21  oct      39         1    151      3
24062 telephone        21  oct     519         1    166      1
24064 telephone        21  oct     144         1     91      4
24072 telephone        22  oct      73         1     86      4
24077 telephone        23  oct     140         1    143      3

      poutcome    y
24060 failure  no
24062 other   yes
24064 failure  yes
24072 other   no
24077 failure  yes
<class 'pandas.core.frame.DataFrame'>
Index: 7842 entries, 24060 to 45210
Data columns (total 17 columns):
 #   Column      Non-Null Count Dtype  
--- 
 0   age          7842 non-null   int64  
 1   job          7842 non-null   object 
 2   marital      7842 non-null   object 
 3   education    7842 non-null   object 
 4   default      7842 non-null   object 
 5   balance      7842 non-null   int64  
 6   housing      7842 non-null   object 
 7   loan          7842 non-null   object 
 8   contact      7842 non-null   object 
 9   day_of_week  7842 non-null   int64  
 10  month         7842 non-null   object 
 11  duration     7842 non-null   int64  
 12  campaign     7842 non-null   int64  
 13  pdays         7842 non-null   int64  
 14  previous     7842 non-null   int64  
 15  poutcome     7842 non-null   object 
 16  y             7842 non-null   object 
dtypes: int64(7), object(10)
memory usage: 1.1+ MB
None

```

```
In [16]: # Target variable: y = "yes" or "no"
df["y"] = df["y"].map({"yes": 1, "no": 0})
df["housing"] = df["housing"].map({"yes": 1, "no": 0})
df["loan"] = df["loan"].map({"yes": 1, "no": 0})
```

Missingness not beyond threshold, Correct datatypes in each column, and No duplicate observations

```
In [17]: # Before validation, check actual missingness
MISSINGNESS_THRESHOLDS = {
    "age": 0.01,
    "job": 0.5,
    "marital": 0.05,
    "education": 0.05,
    "default": 0.05,
    "balance": 0.05,
    "housing": 0.05,
    "loan": 0.05,
    "contact": 0.05,
    "day_of_week": 0.05,
    "month": 0.05,
    "duration": 0.05,
    "campaign": 0.05,
    "pdays": 0.05,
    "previous": 0.05,
    "poutcome": 0.05,
    "y": 0.05,
}

COLUMN_TYPES = {
    "age": pa.Int,
    "job": pa.String,
    "marital": pa.String,
    "education": pa.String,
    "default": pa.Bool,
    "balance": pa.Int,
    "housing": pa.Bool,
    "loan": pa.Bool,
    "contact": pa.String,
    "day_of_week": pa.Int,
    "month": pa.String,
    "duration": pa.Int,
    "campaign": pa.Int,
    "pdays": pa.Int,
    "previous": pa.Int,
    "poutcome": pa.String,
    "y": pa.Bool,
}

for col in df.columns:
    missing_pct = df[col].isna().mean()
    threshold = MISSINGNESS_THRESHOLDS.get(col, 0.05)
    print(f"{col}: {missing_pct:.2%} missing (threshold: {threshold:.2%})")
```

```
age: 0.00% missing (threshold: 1.00%)
job: 0.00% missing (threshold: 50.00%)
marital: 0.00% missing (threshold: 5.00%)
education: 0.00% missing (threshold: 5.00%)
default: 0.00% missing (threshold: 5.00%)
balance: 0.00% missing (threshold: 5.00%)
housing: 0.00% missing (threshold: 5.00%)
loan: 0.00% missing (threshold: 5.00%)
contact: 0.00% missing (threshold: 5.00%)
day_of_week: 0.00% missing (threshold: 5.00%)
month: 0.00% missing (threshold: 5.00%)
duration: 0.00% missing (threshold: 5.00%)
campaign: 0.00% missing (threshold: 5.00%)
pdays: 0.00% missing (threshold: 5.00%)
previous: 0.00% missing (threshold: 5.00%)
poutcome: 0.00% missing (threshold: 5.00%)
y: 0.00% missing (threshold: 5.00%)
```

```
In [18]: import pandera as pa
from pandera import Column, DataFrameSchema, Check

# Helper function to create column schema
def create_column_schema(col_name, col_type, threshold):
    """Create a column schema with missingness check."""
    return Column(
        col_type, # to validates the data type
        checks=[

            Check(
                lambda s: s.isna().mean() <= threshold,
                error=f"Missing value rate exceeds {threshold:.1%}",
                element_wise=False,
                ignore_na=False
            )
        ],
        nullable=True,
        coerce=True # Attempt to coerce to the correct type
    )

# Build schema dynamically
columns_schema = {
    col_name: create_column_schema(col_name, COLUMN_TYPES[col_name], MISSINGNESS_THRESHOLDS.get(col_name, 0.05))
}

schema = DataFrameSchema(
    columns_schema,
    strict=True, # Ensures only defined columns are present
    unique=None, # Will be used for duplicate check below
    coerce=True # Attempt to coerce types across the dataframe
)

# Validate with lazy=True to collect ALL errors
try:
    # First validate schema (types and missingness)
    validated_df = schema.validate(df, lazy=True)
```

```

# Then check for duplicates
duplicates = validated_df.duplicated()
if duplicates.any():
    num_duplicates = duplicates.sum()
    duplicate_rows = validated_df[duplicates]

    print(f"x Found {num_duplicates} duplicate row(s):\n")
    print(duplicate_rows)
    print("\nDuplicate row indices:", duplicate_rows.index.tolist())

    # Optionally raise an error
    raise ValueError(f"DataFrame contains {num_duplicates} duplicate rows")

print("PASSED: All validation checks passed!")
print("- Data types are correct")
print("- Missing value rates are within thresholds")
print("- No duplicate rows found")
print(f"\nValidated DataFrame shape: {validated_df.shape}")
print(validated_df.head())

except pa.errors.SchemaErrors as e:
    print("ERROR: Schema validation failed with the following errors:\n")

    # Display a summary of all failures
    print(f"Total errors found: {len(e.failure_cases)}\n")

    # Check for data type errors
    if 'schema_context' in e.failure_cases.columns:
        type_errors = e.failure_cases[e.failure_cases['check'].str.contains('type')]
        if not type_errors.empty:
            print("Data Type Errors:")
            for _, error in type_errors.iterrows():
                print(f"  • Column '{error['column']}': {error['check']}")
            print()

    # Show which columns failed missingness checks
    failed_columns = e.failure_cases['column'].unique()
    missingness_failures = [col for col in failed_columns
                            if col in MISSINGNESS_THRESHOLDS]

    if missingness_failures:
        print(f"Columns with missingness errors: {missingness_failures}\n")

        # Show detailed error information for each column
        for col in missingness_failures:
            col_errors = e.failure_cases[e.failure_cases['column'] == col]
            actual_missing = df[col].isna().mean()
            threshold = MISSINGNESS_THRESHOLDS.get(col, 0.05)
            print(f"  • {col}:")
            print(f"    - Actual missing rate: {actual_missing:.2%}")
            print(f"    - Threshold: {threshold:.2%}")
            print()

    # Show full error details
    print("\nFull error report:")
    print(e.failure_cases[['schema_context', 'column', 'check', 'check_number']])

```

```

except ValueError as e:
    # Catch duplicate row errors
    print(f"\nx {e}")

except Exception as e:
    print(f"\nx Unexpected error: {e}")

```

PASSED: All validation checks passed!

- Data types are correct
- Missing value rates are within thresholds
- No duplicate rows found

Validated DataFrame shape: (7842, 17)

	age	job	marital	education	default	balance	housing	loan
24060	33	admin.	married	tertiary	True	882	False	False
24062	42	admin.	single	secondary	True	-247	True	True
24064	33	services	married	secondary	True	3444	True	False
24072	36	management	married	tertiary	True	2415	True	False
24077	36	management	married	tertiary	True	0	True	False
24060	telephone		21	oct	39	1	151	3
24062	telephone		21	oct	519	1	166	1
24064	telephone		21	oct	144	1	91	4
24072	telephone		22	oct	73	1	86	4
24077	telephone		23	oct	140	1	143	3
24060	poutcome	y						
24062	failure	False						
24064	other	True						
24072	failure	True						
24077	other	False						
24077	failure	True						

Outlier/Anomaly Validation

In [19]:

```

import pandera as pa
from pandera import Column, Check, DataFrameSchema

```

```

schema = DataFrameSchema(
    {
        "age": Column(
            int,
            checks=[
                Check.in_range(0, 120, include_min=True, include_max=True),
            ],
        ),
        "job": Column(
            str,
            checks=[
                Check(lambda s: s.str.len().between(2, 50, inclusive='both')),
                error="Job title must be between 2 and 50 characters",
            ],
        ),
    }
)

```

```
        ],
        nullable = True
),
"marital": Column(
    str,
    checks=[
        Check.isin(["married", "single", "divorced"])
    ],
),
"education": Column(
    str,
    checks=[
        Check.isin(["tertiary", "secondary", "primary"]),
    ],
    nullable = True
),
"default": Column(
    str,
    checks=[
        Check.isin(["no", "yes"]),
    ],
),
"balance": Column(
    int,
    checks=[
        Check.in_range(-100000, 100000000000, include_min=True, inc
    ],
),
"housing": Column(
    int,
    checks=[
        Check.in_range(0, 1, include_min=True, include_max=True),
    ],
),
"loan": Column(
    int,
    checks=[
        Check.in_range(0, 1, include_min=True, include_max=True),
    ],
),
"contact": Column(
    str,
    checks=[
        Check.isin(["cellular", "telephone"]),
    ],
    nullable = True
),
"day_of_week": Column(
    int,
    checks=[
        Check.in_range(1, 31, include_min=True, include_max=True),
    ],
),
"month": Column(
    str,
    checks=[
```

```

        Check.isin(["jan", "feb", "mar", "apr", "may", "jun", "jul"],
    ],
),
"duration": Column(
    int,
    checks=[
        Check.in_range(0, 21600, include_min=True, include_max=True)
    ],
),
"campaign": Column(
    int,
    checks=[
        Check.in_range(1, 100, include_min=True, include_max=True),
    ],
),
"pdays": Column(
    int,
    checks=[
        Check.in_range(-1, 10000, include_min=True, include_max=True),
    ],
),
"previous": Column(
    int,
    checks=[
        Check.in_range(0, 1000000, include_min=True, include_max=True),
    ],
),
"poutcome": Column(
    str,
    checks=[
        Check.isin(["failure", "other", "success"]),
    ],
    nullable = True
),
"y": Column(
    int,
    checks=[
        Check.in_range(0, 1, include_min=True, include_max=True),
    ],
),
},
strict=True,
coerce=True,
)

def is_outlier_iqr(series, multiplier=1.5):
    Q1 = series.quantile(0.25)
    Q3 = series.quantile(0.75)
    IQR = Q3 - Q1
    lower_bound = Q1 - multiplier * IQR
    upper_bound = Q3 + multiplier * IQR
    return (series < lower_bound) | (series > upper_bound)

try:
    validated_df = schema.validate(df)
    print("✓ Data validation passed!")

```

```

print(validated_df)

except pa.errors.SchemaError as e:
    print("x Data validation failed!")
    print(f"\nError summary: {e}")
    print("\nDetailed failure cases:")
    print(e.failure_cases)

✓ Data validation passed!
   age      job marital education default balance housing loan
\ 
24060  33    admin. married  tertiary no     882     0     0
24062  42    admin. single secondary no    -247     1     1
24064  33    services married secondary no    3444    1     0
24072  36    management married tertiary no    2415    1     0
24077  36    management married tertiary no     0     1     0
...
45199  34  blue-collar single secondary no    1475    1     0
45201  53    management married tertiary no     583     0     0
45204  73    retired married secondary no    2850     0     0
45208  72    retired married secondary no    5715     0     0
45210  37 entrepreneur married secondary no    2971     0     0

   contact day_of_week month duration campaign pdays previous \
24060 telephone          21  oct       39        1    151      3
24062 telephone          21  oct      519        1    166      1
24064 telephone          21  oct      144        1     91      4
24072 telephone          22  oct       73        1     86      4
24077 telephone          23  oct      140        1    143      3
...
45199 cellular           16 nov      1166       3    530     12
45201 cellular           17 nov      226        1    184      4
45204 cellular           17 nov      300        1     40      8
45208 cellular           17 nov     1127       5    184      3
45210 cellular           17 nov      361       2    188     11

   poutcome y
24060 failure 0
24062 other 1
24064 failure 1
24072 other 0
24077 failure 1
...
45199 other 0
45201 success 1
45204 failure 1
45208 success 1
45210 other 0

[7842 rows x 17 columns]

```

Correct category levels (i.e., no string mismatches or single values)

```
In [20]: categorical_df = df[["job", "marital", "education", "contact", "month", "pou  
schema = DataFrameSchema(  
    {  
        "job": Column(  
            str,  
            checks=[  
                Check(lambda s: s.str.len().between(2, 50, inclusive='both'))  
                .error="Job title must be between 2 and 50 characters"),  
            nullable = True  
        ),  
        "marital": Column(  
            str,  
            checks=[  
                Check.isin(["married", "single", "divorced"])  
            ],  
        ),  
        "education": Column(  
            str,  
            checks=[  
                Check.isin(["tertiary", "secondary", "primary"]),  
            ],  
            nullable = True  
        ),  
        "contact": Column(  
            str,  
            checks=[  
                Check.isin(["cellular", "telephone"]),  
            ],  
            nullable = True  
        ),  
        "month": Column(  
            str,  
            checks=[  
                Check.isin(["jan", "feb", "mar", "apr", "may", "jun", "jul",  
            ],  
        ),  
        "poutcome": Column(  
            str,  
            checks=[  
                Check.isin(["failure", "other", "success"]),  
            ],  
            nullable = True  
        ),  
    },  
    strict=True,  
    coerce=True,  
)  
  
try:  
    validated_df = schema.validate(categorical_df)  
    print("\u2713 Data validation passed!")  
    print(validated_df)
```

```

except pa.errors.SchemaError as e:
    print("x Data validation failed!")
    print(f"\nError summary: {e}")
    print("\nDetailed failure cases:")
    print(e.failure_cases)

```

✓ Data validation passed!

	job	marital	education	contact	month	poutcome
24060	admin.	married	tertiary	telephone	oct	failure
24062	admin.	single	secondary	telephone	oct	other
24064	services	married	secondary	telephone	oct	failure
24072	management	married	tertiary	telephone	oct	other
24077	management	married	tertiary	telephone	oct	failure
...
45199	blue-collar	single	secondary	cellular	nov	other
45201	management	married	tertiary	cellular	nov	success
45204	retired	married	secondary	cellular	nov	failure
45208	retired	married	secondary	cellular	nov	success
45210	entrepreneur	married	secondary	cellular	nov	other

[7842 rows x 6 columns]

In [21]: # 4. Split features and target

```

# Identify numerical and categorical columns
numerical_cols = X.select_dtypes(include=["int64", "float64"]).columns
categorical_cols = X.select_dtypes(include=["object"]).columns

print(numerical_cols)
print(categorical_cols)

Index(['age', 'balance', 'day_of_week', 'duration', 'campaign', 'pdays',
       'previous'],
      dtype='object')
Index(['job', 'marital', 'education', 'default', 'housing', 'loan', 'contact',
       'month', 'poutcome'],
      dtype='object')

```

```
In [22]: # 5. Preprocessing pipeline
numeric_transformer = Pipeline(
    steps=[
        ("scaler", StandardScaler())
    ])

categorical_transformer = Pipeline(
    steps=[
        ("onehot", OneHotEncoder(handle_unknown="ignore"))
    ])

preprocessor = ColumnTransformer(
    transformers=[
        ("num", numeric_transformer, numerical_cols),
        ("cat", categorical_transformer, categorical_cols)
    ]
)
```

Fitting the model and making predictions

```
In [23]: # 6. Build model pipeline
model = Pipeline(steps=[
    ("preprocessor", preprocessor),
    ("classifier", LogisticRegression(max_iter=1000, class_weight ="balanced"))
])
```

```
In [24]: # 7. Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

Data Validation - Target/Response Variable Follows Expected Distribution

```
In [25]: from deepchecks.tabular.checks import LabelDrift , TrainTestFeatureDrift, Mu
from deepchecks.tabular import Dataset, Suite

categorical_feats = ["job", "marital", "education", "contact", "month", "po
training_dataset = Dataset(pd.concat([X_train, y_train]), label = "y", cat_f
testing_dataset = Dataset(pd.concat([X_test, y_test]), label = "y", cat_feat

distribution_check = Suite(
    'Drift Detection Suite',
    LabelDrift().add_condition_drift_score_less_than(0.15),
    MultivariateDrift().add_condition_overall_drift_value_less_than(0.15)
)

suite_result = distribution_check.run(training_dataset, testing_dataset)
has_drift = False
for check_result in suite_result.results:
```

```
print(f"\nCheck: {check_result.get_header()}\")\n\n    if check_result.conditions_results:\n        for condition_result in check_result.conditions_results:\n            status = "✓" if condition_result.is_pass() else "✗"\n            print(f"  {status} {condition_result.name}")\n            if not condition_result.is_pass():\n                has_drift = True\n                print(f"      Details: {condition_result.details}")\n    else:\n        print("  No conditions set for this check")
```

/Users/rabin/miniforge3/envs/522/lib/python3.12/site-packages/deepchecks/tabular/dataset.py:236: UserWarning:

Dataframe index has duplicate indexes, setting index to [0,1..,n-1].

/Users/rabin/miniforge3/envs/522/lib/python3.12/site-packages/deepchecks/tabular/dataset.py:236: UserWarning:

Dataframe index has duplicate indexes, setting index to [0,1..,n-1].

```
/Users/rabin/miniforge3/envs/522/lib/python3.12/multiprocessing/queues.py:12
  2: UserWarning: pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated for removal as early as 2025-11-30. Refrain from using this package or pin to Setuptools<81.
      return _ForkingPickler.loads(res)
/Users/rabin/miniforge3/envs/522/lib/python3.12/multiprocessing/queues.py:12
  2: UserWarning: pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated for removal as early as 2025-11-30. Refrain from using this package or pin to Setuptools<81.
      return _ForkingPickler.loads(res)
/Users/rabin/miniforge3/envs/522/lib/python3.12/multiprocessing/queues.py:12
  2: UserWarning: pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated for removal as early as 2025-11-30. Refrain from using this package or pin to Setuptools<81.
      return _ForkingPickler.loads(res)
/Users/rabin/miniforge3/envs/522/lib/python3.12/multiprocessing/queues.py:12
  2: UserWarning: pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated for removal as early as 2025-11-30. Refrain from using this package or pin to Setuptools<81.
      return _ForkingPickler.loads(res)
/Users/rabin/miniforge3/envs/522/lib/python3.12/multiprocessing/queues.py:12
  2: UserWarning: pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated for removal as early as 2025-11-30. Refrain from using this package or pin to Setuptools<81.
      return _ForkingPickler.loads(res)
/Users/rabin/miniforge3/envs/522/lib/python3.12/multiprocessing/queues.py:12
  2: UserWarning: pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated for removal as early as 2025-11-30. Refrain from using this package or pin to Setuptools<81.
      return _ForkingPickler.loads(res)
/Users/rabin/miniforge3/envs/522/lib/python3.12/multiprocessing/queues.py:12
  2: UserWarning: pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated for removal as early as 2025-11-30. Refrain from using this package or pin to Setuptools<81.
      return _ForkingPickler.loads(res)
/Users/rabin/miniforge3/envs/522/lib/python3.12/multiprocessing/queues.py:12
  2: UserWarning: pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated for removal as early as 2025-11-30. Refrain from using this package or pin to Setuptools<81.
      return _ForkingPickler.loads(res)
/Users/rabin/miniforge3/envs/522/lib/python3.12/multiprocessing/queues.py:12
  2: UserWarning: pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated for removal as early as 2025-11-30. Refrain from using this package or pin to Setuptools<81.
```

```
s.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated  
for removal as early as 2025-11-30. Refrain from using this package or pin t  
o Setuptools<81.  
    return _ForkingPickler.loads(res)  
/Users/rabin/miniforge3/envs/522/lib/python3.12/multiprocessing/queues.py:12  
2: UserWarning: pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated  
for removal as early as 2025-11-30. Refrain from using this package or pin t  
o Setuptools<81.  
    return _ForkingPickler.loads(res)  
/Users/rabin/miniforge3/envs/522/lib/python3.12/multiprocessing/queues.py:12  
2: UserWarning: pkg_resources is deprecated as an API. See https://setuptools.pypa.io/en/latest/pkg_resources.html. The pkg_resources package is slated  
for removal as early as 2025-11-30. Refrain from using this package or pin t  
o Setuptools<81.  
    return _ForkingPickler.loads(res)  
Check: Label Drift  
    ✓ Label drift score < 0.15
```


```
Check: Multivariate Drift  
    ✓ Drift value is less than 0.15
```

```
In [26]: # 8. Train model  
model.fit(X_train, y_train)
```

```
/Users/rabin/miniforge3/envs/522/lib/python3.12/site-packages/scikit-learn/utils/  
validation.py:1406: DataConversionWarning:
```

A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n_samples,), for example using ravel().

```
Out[26]: ► Pipeline
```



```
In [27]: # 9. Predictions and evaluation
```

```
y_pred = model.predict(X_test)  
y_prob = model.predict_proba(X_test)[:, 1]  
  
print("Accuracy:", accuracy_score(y_test, y_pred))  
print("ROC-AUC:", roc_auc_score(y_test, y_prob))  
print("\nClassification Report:\n", classification_report(y_test, y_pred))  
  
# Confusion matrix  
cm = confusion_matrix(y_test, y_pred)
```

```

sns.heatmap(cm,
            annot=True,
            fmt="d",
            cmap="Blues",
            xticklabels=["no", "yes"],
            yticklabels=["no", "yes"]
)
plt.title("Confusion Matrix")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.show()


```

Accuracy: 0.8457370341700763

ROC-AUC: 0.9079218714674134

Classification Report:

	precision	recall	f1-score	support
no	0.97	0.85	0.91	7985
yes	0.42	0.81	0.55	1058
accuracy			0.85	9043
macro avg	0.70	0.83	0.73	9043
weighted avg	0.91	0.85	0.87	9043



In [28]: # 10. Feature Importance (for logistic regression)
This is a bit tricky with pipelines – we extract processed feature names
ohe = model.named_steps["preprocessor"].named_transformers_["cat"]["onehot"]
cat_feature_names = ohe.get_feature_names_out(categorical_cols)

```

feature_names = np.concatenate([numerical_cols, cat_feature_names])


# Get coefficients
coeffs = model.named_steps["classifier"].coef_[0]

feat_imp = pd.DataFrame({
    "feature": feature_names,
    "importance": coeffs
}).sort_values(by="importance", ascending=False)

print(feat_imp.head(10))
feat_imp.head(20).plot(kind="bar", x="feature", y="importance", figsize=(10,
plt.title("Feature Importance (Logistic Regression Coefficients)")
plt.show()

```

	feature	importance
49	poutcome_success	1.824477
42	month_mar	1.722014
3	duration	1.525751
45	month_oct	1.276304
46	month_sep	0.984794
37	month_dec	0.713720
15	job_student	0.657311
12	job_retired	0.556467
33	contact_telephone	0.546549
32	contact_cellular	0.522132



References

A data modeling approach for classification problems: application to bank telemarketing prediction Stéphane Cédric KOUMETIO TEKOUABOU, Walid Cherif, H. Silkan Published in International Conferences on... 27 March 2019 Computer Science, Business <https://www.semanticscholar.org/paper/A-data-modeling-approach-for-classification-to-bank-TEKOUABOU-Cherif/241d6ca92c4bc65ac3ee903e4732f70bff5c5e9f>

Predicting the Accuracy for Telemarketing Process in Banks Using Data Mining F. Alsolami, Farrukh Saleem, A. Al-malaise, AL-Ghamdi, Published 2020 Business, Computer Science <https://www.semanticscholar.org/paper/Predicting-the-Accuracy-for-Telemarketing-Process-Alsolami-Saleem/6391b7edcdd3c443bb57624b153bf9a8cca027db>

Using Logistic Regression Model to Predict the Success of Bank Telemarketing Y. Jiang, Published 21 June 2018 Business, Computer Science, Journal of data science <https://www.semanticscholar.org/paper/Using-Logistic-Regression-Model-to-Predict-the-of-Jiang/11ea58c843d0e745716d624b03067235dc285c30>

Prediction of Term Deposit in Bank: Using Logistic Model Enjing Jiang, Zihao Wang, Jiaying Zhao, Published in BCP Business & Management 14 December 2022 Business, Computer Science <https://www.semanticscholar.org/paper/Prediction-of-Term-Deposit-in-Bank%3A-Using-Logistic-Jiang-Wang/e36cafceaad636e9b2b558166c16be31a913ad0d>